

Information Storage and Retrieval

Tuesday, December 09, 2025 3:42 PM

Unit 06: Advanced Information Retrieval

Nov – Dec 2022

Q7) a) Define Recommender system? Explain in brief Collaborative Filtering.

Ans.

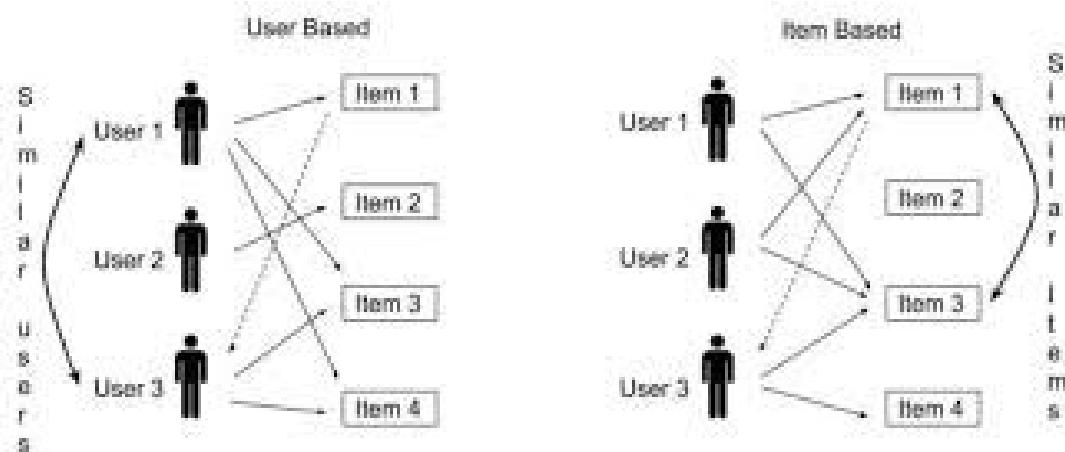
What is a Recommender System?

A **Recommender System** is an information filtering tool that suggests items to users based on their preferences, behavior, or similarities with other users/items.

It is widely used in platforms like **Netflix (movie recommendations)**, **Amazon (product suggestions)**, **Spotify (music recommendations)**, etc.

Purpose of a Recommender System

- To help users discover relevant items from a large collection.
- To improve user satisfaction and engagement.
- To personalize user experience.



What is Collaborative Filtering?

Collaborative Filtering (CF) is one of the most common approaches used in recommender systems.

It makes predictions about a user's interests by **collecting preferences from many other users**.

The key idea is:

"Users who agreed in the past tend to agree in the future."

Types of Collaborative Filtering

1. User-Based Collaborative Filtering

- Finds users similar to the target user.

- Recommends items those similar users liked.
- Example: If User A and User B like the same movies, and User B likes another movie X, then recommend X to User A.

2. Item-Based Collaborative Filtering

- Finds similarity between items.
- Recommends items similar to the ones the user already likes.
- Example: If many users who watched Movie A also watched Movie B, then Movie B is recommended to a user who watched A.

How Collaborative Filtering Works (Steps)

1. **Collect user-item interaction data**
 - Ratings, clicks, purchases, watch history.
2. **Calculate similarity**
 - Between users (user-based) or items (item-based).
 - Using cosine similarity, Pearson correlation, etc.
3. **Predict rating or preference**
 - Based on similar users/items.
4. **Recommend top items**
 - Highest predicted ratings or relevance.

Advantages of Collaborative Filtering

- Does not require item information/content.
- Learns complex user preferences automatically.
- Provides personalized recommendations.

Limitations

- **Cold Start Problem**
New users/items have little data, making recommendations hard.
- **Sparsity**
Most users rate only a few items, causing sparse matrices.
- **Scalability**
Large datasets require heavy computation.

Q7) b) Explain semantic web in details.

Ans.

What is the Semantic Web?

The **Semantic Web** is an extension of the current World Wide Web that aims to make data **understandable by machines**.

Instead of just displaying information for humans to read, the Semantic Web allows computers to **interpret, share, and connect** data automatically.

It was proposed by **Tim Berners-Lee**, the inventor of the Web.

Key Idea

"Give meaning (semantics) to web data so machines can understand relationships between things."

This transforms the web from a collection of documents into a **web of interconnected data**.

Why do we need the Semantic Web?

- People understand meaning, but machines do not.
- Websites use different formats and structures.
- With semantics, machines can **combine information** from many sources.

Example:

A machine could understand that:

- "Paris is the capital of France"
 - "France is in Europe"
- So it can infer: "Paris is a city in Europe."

How the Semantic Web Works

It uses technologies that add **metadata** (data about data) to web pages so machines can understand context.

Main Components

5. RDF (Resource Description Framework)

- A model to represent information in triples:
Subject – Predicate – Object

Example:

Paris – isCapitalOf – France

6. OWL (Web Ontology Language)

- Defines concepts, relationships, and rules (like a dictionary for machines).

7. SPARQL

- A query language used to retrieve and manipulate semantic data.

8. URLs (Unique Identifiers)

- Identify resources uniquely on the web.

Semantic Web Architecture (Layer Cake)





This diagram shows the stack of technologies used in the Semantic Web, from basic data encoding to reasoning and trust layers.

Benefits of the Semantic Web

- Better search results (understanding meaning, not just keywords)
- Data integration across multiple platforms
- Intelligent applications (AI systems)
- Automated reasoning and inference
- More meaningful and connected information

Example

Without Semantic Web:

The word “Apple” might mean the fruit or the company—machines cannot tell.

With Semantic Web:

Metadata clarifies:

- Apple (Company) → Technology organization
- Apple (Fruit) → Edible fruit

Machines can differentiate and link related knowledge.

In Short

The Semantic Web turns the web into a **machine-understandable knowledge network**, enabling smarter applications, better search, and interconnected data.

Q8) a) Explain difference between Text-centric and Data-centric XML retrieval.

Ans.

Comparison Table: Text-Centric vs Data-Centric XML Retrieval

Feature	Text-Centric XML Retrieval	Data-Centric XML Retrieval
Content Type	Long, narrative, unstructured or semi-structured text	Highly structured, well-defined data
Purpose	Document organization and content markup	Data storage, exchange, and processing
Structure Strictness	Loose structure, flexible	Rigid, schema-driven

Typical Data	Articles, reports, books, blogs	Customer records, product catalogs, transactions
Element Size	Large elements (paragraphs, sections)	Small, atomic data fields (name, price, ID)
Primary Retrieval Goal	Find relevant text sections	Retrieve specific data values
Query Type	Keyword search, full-text search, ranking	Structured queries (XPath, XQuery)
Processing Style	Similar to Information Retrieval (IR)	Similar to Database Management Systems (DBMS)
Tag Function	Helps structure narrative content	Defines data model and relationships
Examples of Use	Digital libraries, news search, document databases	E-commerce databases, banking data, inventory systems

Q8) b) Explain in detail Content Based Recommendation of Documents.

Ans.

What is Content-Based Recommendation of Documents?

Content-Based Recommendation is a method of suggesting documents to a user **based on the content (features)** of documents the user has shown interest in.

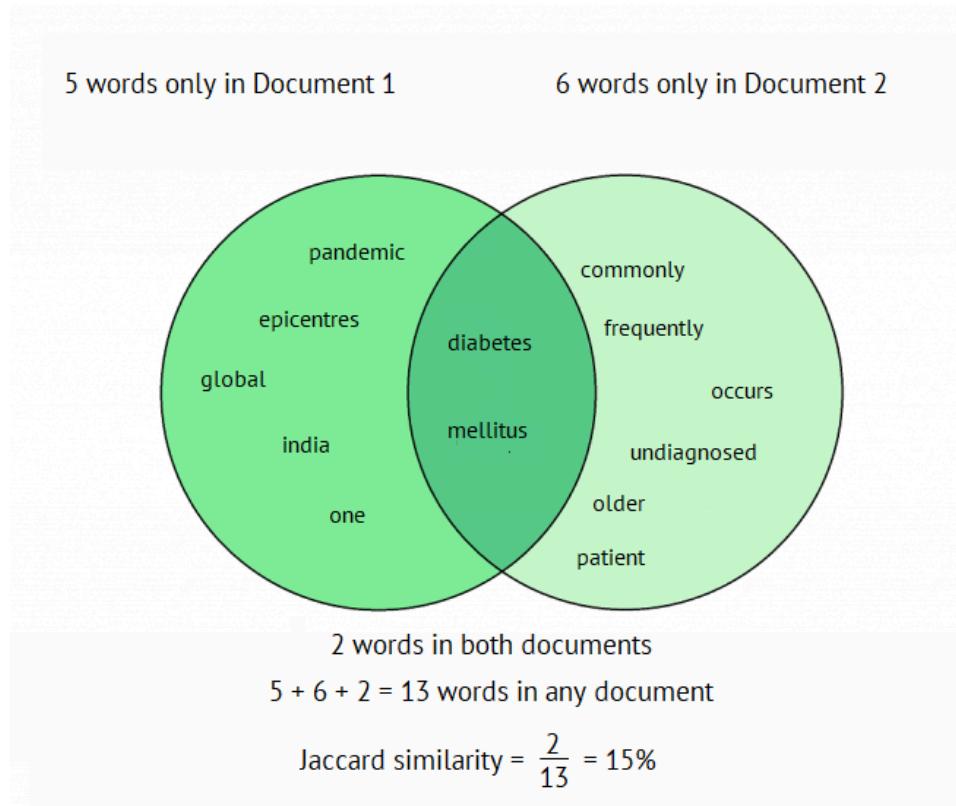
The main idea:

"Recommend items similar to those the user already likes, based on document content."

This approach does **not** depend on other users' behavior (unlike Collaborative Filtering).

It analyzes the *features* of documents such as keywords, topics, metadata, and relevance.

How It Works (Step-by-Step)



1. Represent Document Content

Documents are converted into a structured form using features such as:

- Keywords
- Term Frequency – Inverse Document Frequency (TF-IDF)
- Bag-of-Words (BoW)
- LDA topic modeling
- Word embeddings (Word2Vec, BERT)

2. Build User Profile

The system creates a profile of user preferences based on:

- Documents the user clicked, read, rated, or liked
- The frequent terms/topics in these documents

A user profile might include:

- Preferred topics
- Top keywords
- Semantic vectors

3. Compute Similarity

The system measures similarity between:

- The user profile
and
- Other documents in the collection

Common similarity measures:

- Cosine similarity
- Jaccard similarity
- Euclidean distance

4. Recommend Most Similar Documents

Documents with the highest similarity score are recommended to the user.

Example

Suppose a user frequently reads documents related to:

- “Machine Learning”
- “Neural Networks”
- “AI Models”

The system identifies these topics in their profile and recommends documents containing similar terms or topics, such as:

- “Deep Learning Overview”
- “Training Neural Networks”
- “Applications of AI in Healthcare”

Advantages of Content-Based Recommendation

- **No cold-start for items** (new documents can be recommended immediately)
- **Personalized recommendations** based on user's own behavior
- **No need for other user data**
- **Explainable** ("Recommended because it contains similar topics as your previous reads")

Limitations

- **Limited novelty/diversity**

Recommends items too similar to what the user already knows.

- **Requires good feature extraction**

Poorly represented documents → poor recommendations.

- **Cold-start for new users**

Needs some user history to build a profile.

In Summary

Content-Based Document Recommendation uses document features and user preferences to find and recommend similar documents. It works by analyzing content, building a user profile, and computing similarity between documents.

May – Jun 2023

Q7) a) Differentiate Collaborative filtering and Content Filtering.

Ans.

Comparison Table

Feature	Collaborative Filtering (CF)	Content-Based Filtering (CBF)
Basis of Recommendation	Uses preferences/behaviors of <i>similar users</i>	Uses <i>content/features</i> of items the user liked
Requires Other Users' Data?	Yes	No
Key Idea	"People similar to you liked this."	"You liked this item; here are similar ones."
Data Used	Ratings, clicks, interactions from many users	Item attributes (keywords, genres, topics, metadata)
Cold Start Issues	Struggles with new users/items (lack of data)	New items can be recommended easily; new users still a problem
Novelty / Diversity	Can recommend unexpected items based on community trends	Tends to recommend similar items → limited novelty
Scalability	Harder with huge user–item matrices (sparse data)	More scalable because it uses item features

Explainability	Harder to explain (“similar users liked it”)	Easy to explain (“shares features with items you liked”)
Example Scenario	Netflix recommends a movie because many users with similar taste watched it	Spotify recommends songs similar to the one you played
Dependency on Item Features	Not required	Essential (needs good feature extraction)

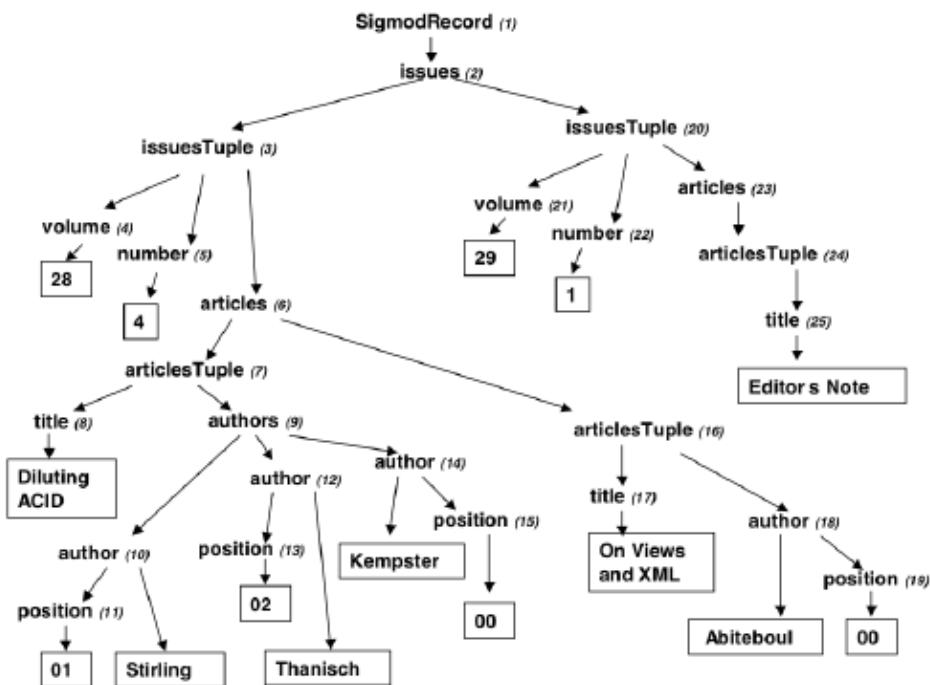
Simple One-Line Difference

- **Collaborative Filtering:** Recommends items based on similar users' preferences.
- **Content-Based Filtering:** Recommends items similar to those the user already likes.

Q8) a) Explain Text-Centric and Data-Centric XML retrieval.

Ans.

1. Text-Centric XML Retrieval



Definition

Text-centric XML retrieval deals with XML documents where the **primary content is text**—often long, narrative, or semi-structured.

XML tags are used mainly for **marking sections**, not for enforcing strict data structures.

Characteristics

- Contains **large blocks of human-readable text** (paragraphs, articles, reports).
- Structure is **loose and flexible**.
- Retrieval resembles **Information Retrieval (IR)**.
- Uses ranking techniques like keyword matching, TF-IDF, etc.

Examples

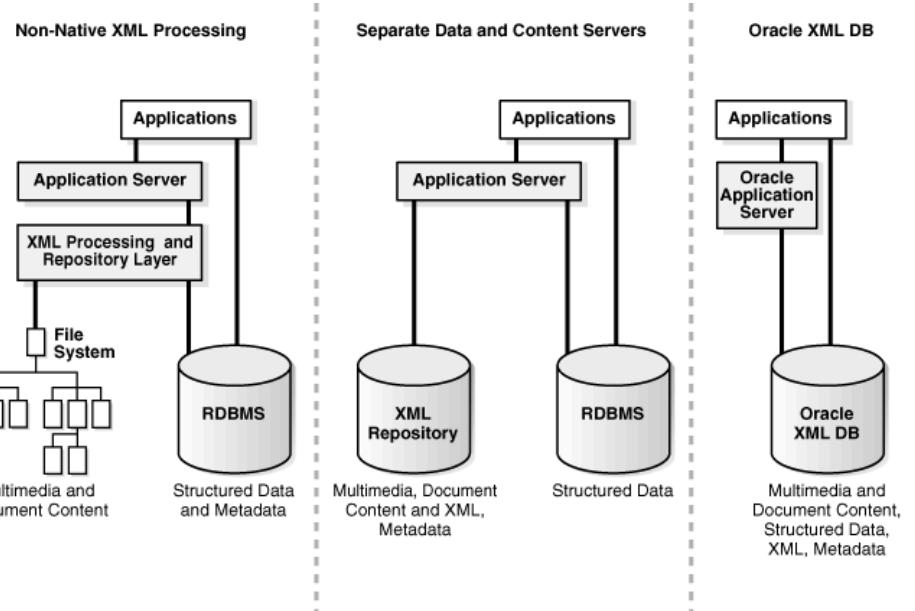
- News articles
- Research papers

- Blogs
- Digital library content

Retrieval Goal

- Find **relevant text fragments** or documents based on keywords or topics.

2. Data-Centric XML Retrieval



Definition

Data-centric XML retrieval involves XML documents where information is **highly structured**, similar to database records.

XML is used mainly to store and exchange structured data.

Characteristics

- Contains **small, well-defined fields** (numbers, names, dates).
- Structure is **strict and follows a schema**.
- Retrieval similar to **database querying**.
- Uses XPath, XQuery, and exact matching.

Examples

- Banking transactions
- Product catalogs
- Customer databases
- Inventory records

Retrieval Goal

- Retrieve **specific data values** efficiently and accurately.

3. Key Differences (Summary Table)

Feature	Text-Centric XML Retrieval	Data-Centric XML Retrieval
Content Type	Narrative text	Structured data
Structure	Loose, semi-structured	Rigid, schema-driven
Typical Elements	Paragraphs, sections	Small atomic data fields

Query Type	Keyword-based, ranked retrieval	Structured queries (XPath/XQuery)
Similar To	Search engines	Databases
Retrieval Output	Text fragments, whole documents	Exact records/values
Application	Digital libraries	Data storage/exchange systems

In Simple Words

- **Text-Centric XML** → Used for documents with **lots of text**; retrieved like search engines.
- **Data-Centric XML** → Used for **structured data**; retrieved like database queries.

Nov– Dec 2023

Q7) b) Explain Vector Space Model for XML Retrieval.

Ans. Below is a **clear explanation of the Vector Space Model (VSM) for XML Retrieval**, along with a helpful diagram-style image.

Vector Space Model (VSM) for XML Retrieval

The **Vector Space Model** is a mathematical model used in Information Retrieval (IR) to represent documents and queries as **vectors in a multi-dimensional space**.

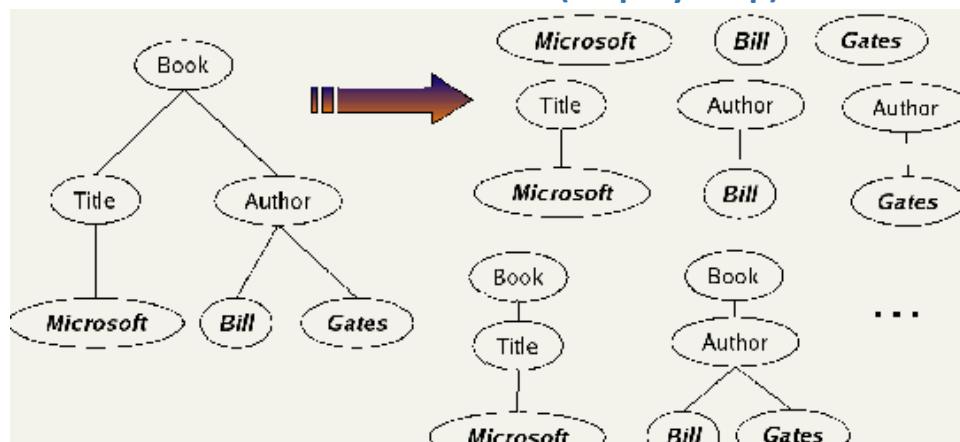
For **XML retrieval**, VSM is adapted to work not just with whole documents but also with **XML elements**, since users may want to retrieve specific sections of an XML document.

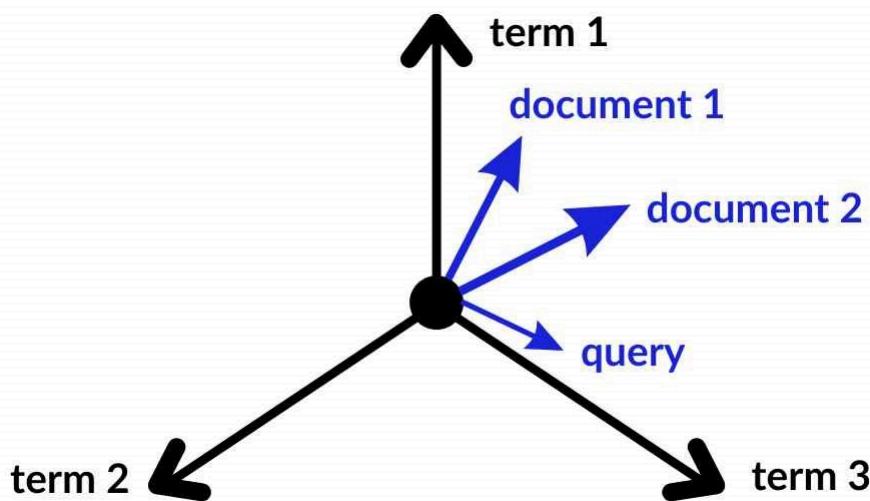
Key Idea

Represent XML elements and queries as vectors, compute similarity, and rank the most relevant XML elements.

Each dimension represents a **term**, and each XML element is represented as a vector of term weights (often TF-IDF).

How VSM Works for XML Retrieval (Step-by-Step)





1. Indexing XML Documents

XML documents are broken down into:

- Whole documents
- Elements (e.g., `<section>`, `<paragraph>`, `<article>`, etc.)

Each element is treated as a separate retrievable unit.

2. Representing XML Elements as Vectors

For each XML element:

- Extract terms (after stemming, stopword removal)
- Compute term weights (TF-IDF)
- Create a vector such as:
→ $d = (w_1, w_2, w_3, \dots, w_n)$

where (w_i) is the weight of term (i).

3. Representing the Query as a Vector

A query like:

“climate change impact”

is represented as another vector with term weights.

4. Calculating Similarity (Cosine Similarity)

To rank XML elements, compute:

```
[  
text{Similarity}(d,q)=\frac{\vec{d} \cdot \vec{q}}{\|\vec{d}\| \|\vec{q}\|}  
]
```

Higher cosine value → more relevant element.

5. Ranking and Returning Results

The system returns:

- The most relevant **XML fragments**,
not necessarily whole documents.

This makes VSM ideal for **text-centric XML**.

Why Use VSM in XML Retrieval?

Advantages

- Supports retrieval of **specific XML parts** (granularity).
- Produces **ranked results** based on similarity.
- Works well for **text-centric XML** (articles, reports).
- Easy to implement with TF-IDF.

Limitations

- Struggles with **data-centric XML** (numeric/structured data).
- Ignores XML structure unless structural weighting is added.
- Result quality depends on term weighting.

Example

Suppose an XML article has sections on:

- Climate change
- Weather patterns
- Environmental impacts

A query “climate impact” will return the **specific XML element(s)** most similar to the query vector, not just the whole document.

In Simple Terms

The Vector Space Model for XML retrieval:

- Converts XML elements into vectors
- Converts the user query into a vector
- Measures similarity
- Returns the most relevant XML fragment

It is a powerful method for retrieving **textual XML content**.

Nov– Dec 2024

Q7) b) Write short notes on Challenges in XML Retrieval.

Ans.

Challenges in XML Retrieval

XML Retrieval is more complex than traditional document retrieval because XML contains **hierarchical structure**, **nested elements**, and **mixed content**. Below are the major challenges:

1. Structural Complexity

XML documents are organized hierarchically (elements inside elements).

Different documents may use different structures even when representing similar content.

Challenge:

- Understanding and matching the structure during retrieval is difficult.
- Similar content may appear at different levels of the hierarchy.

2. Granularity of Retrieval

XML allows retrieving:

- Whole documents
- Sections
- Paragraphs
- Small elements

Challenge:

- Determining the correct retrieval unit (how big or small the answer should be).
- Ranking results with different sizes is difficult.

Example:

Which should rank higher: a whole article or a single relevant paragraph?

3. Mixed Content (Text + Structure)

XML documents contain both **textual content** and **tags**.

Challenge:

- How much importance to give to text vs. structure?
- Structural tags may or may not be relevant to the user's query.

4. Heterogeneous Schemas

Different publishers or systems may use different tag names or structures for similar data.

Challenge:

- Hard to match queries across multiple XML formats.
- Requires schema mapping or ontology knowledge.

Example:

One document uses `<author>`, another uses `<writer>` for the same concept.

5. Ranking XML Elements

Traditional ranking (like VSM or TF-IDF) works on full documents, but XML returns *fragments*.

Challenge:

- How to rank elements of different size and depth?
- Relevance propagation between parent and child elements is complex.

6. Overlapping Retrieval Results

Since XML is nested:

- A paragraph may be relevant
- Its section may also be relevant
- The entire article may also be relevant

Challenge:

- Avoiding duplicate or overlapping results.
- Deciding which level should be shown to the user.

7. Query Ambiguity (Content vs Structure)

XML queries may specify:

- Only text terms
- Only structure
- Both (content + structure)

Challenge:

Interpreting complex queries like:

Find `<section>` elements containing "climate change"

requires understanding both textual and structural constraints.

8. Efficient Indexing

XML has many small elements → millions of retrievable units.

Challenge:

- Indexing every element is expensive.
- Need special indexes for structure-aware searching (path indexes, element indexes).

9. Performance and Scalability

XML repositories can be very large, with deeply nested structures.

Challenge:

- Complex structural matching slows search.
- Memory and storage overhead is higher compared to flat documents.

10. Evaluation Difficulty

Traditional IR metrics (precision, recall) must be adapted.

Challenge:

- How to evaluate partial or nested results?

- Multiple correct answers may exist at different granularities.

Summary

XML Retrieval is challenging because XML adds **structure**, **granularity**, and **hierarchy** to traditional text retrieval.

Systems must deal with:

- Variable structures
- Element-level retrieval
- Ranking and overlap reduction
- Schema differences
- Efficient indexing

Basic XML Concepts

XML (**Extensible Markup Language**) is a markup language designed to store and transport data in a **structured, self-descriptive** format. It is widely used for data exchange in web applications, databases, and documents.

1. XML Document Structure

An XML document consists of **elements**, **attributes**, **tags**, and **data**.

Example XML

```
<book> <title>XML Basics</title> <author>John Doe</author> <price
currency="USD">29.99</price> </book>
```

2. Elements

Elements are the building blocks of XML.

They contain:

- Opening tag
- Content
- Closing tag

Example:

```
<title>XML Basics</title>
```

Elements can also contain other elements (nested structure).

3. Attributes

Attributes provide **additional information** about elements.

Example:

```
<price currency="USD">29.99</price>
```

Here, `currency="USD"` is an attribute of the `price` element.

Note:

Attributes store metadata; elements store actual data.

4. Tags

Tags define the start and end of an XML element.

- Opening tag: <author>
- Closing tag: </author>
- Empty tag: <linebreak/>

Tags must be:

- Properly nested
- Case-sensitive

5. Hierarchical (Tree) Structure

XML organizes data in a **tree structure** with:

- A **root element** (only one)
- Child elements
- Parent-child relationships
- Sibling elements

Example:

```
<library> < Root <book> < Child </book> </library>
```

6. XML Declaration

Optional line at the beginning of XML:

```
<?xml version="1.0" encoding="UTF-8"?>
```

Specifies XML version and character encoding.

7. Well-Formed XML

An XML document is “well-formed” if it follows syntax rules:

- One root element
- Properly nested tags
- Closing tags present
- Case-sensitive tags
- Attribute values in quotes

Example (valid):

```
<note> <to>Alice</to> <from>Bob</from> </note>
```

8. Valid XML

An XML document is “valid” if it follows rules defined in:

- **DTD** (Document Type Definition), or
- **XML Schema (XSD)**

Validation ensures structure and data types are correct.

9. XML Namespaces

Namespaces avoid naming conflicts when combining XML from different sources.

Example:

```
<bookstore xmlns:fiction="http://example.com/fiction"> <fiction:book>The  
Alchemist</fiction:book> </bookstore>
```

10. XML Comments

Comments are written as:

```
<!-- This is a comment -->
```

11. Character Data (CDATA)

CDATA sections store text that should not be parsed by XML processors.

Example:

```
<![CDATA[ <h1>This is not XML code</h1> ]]>
```

12. XML Advantages

- Human-readable and machine-readable
- Platform-independent
- Supports hierarchical data
- Good for data exchange
- Extensible (user-defined tags)

In Summary

XML is a flexible, self-descriptive language used to represent structured data.

Understanding elements, attributes, tags, structure, and validation concepts is essential for working with XML in databases, web services, and information retrieval.